

Enhancing Efficiency of Korean Chatbots with a DynamicTanh Transformer Model

Cheol Hun Yeom*
AIFEL Research Generation 12
drhunny1@gmail.com

March 23, 2025

Abstract

This study investigates the effectiveness of incorporating a DynamicTanh (DyT) layer (Zhu *et al.*, 2025), used as a replacement for LayerNorm (Ba *et al.*, 2016), within a Transformer model for constructing a Korean-language chatbot. Leveraging their self-attention mechanism, Transformers have demonstrated proficiency in handling long-range dependencies within text, leading to strong performance in various chatbot applications. We train a Transformer-based chatbot on a Korean question-answer dataset and conduct a comparative analysis between a model employing DyT and a baseline (LayerNorm) model. Performance is primarily evaluated using the BLEU (Bilingual Evaluation Understudy) score (Papineni *et al.*, 2002). Our results reveal that both the DyT and LayerNorm models achieved an average BLEU score of approximately 64 on the training data (n=100). However, the DyT model achieved a speedup of approximately 15% in training and evaluation time compared to the LayerNorm model. These observations suggest that the DyT layer can serve as a promising replacement for LayerNorm, enhancing efficiency in Transformer-based Korean chatbots without sacrificing performance on the training data.

1 Introduction

Chatbots are becoming increasingly common in many areas, like customer service and providing information, thanks to advances in Natural Language Processing (NLP). Building chatbots for the Korean language is especially challenging because of its complex grammar and structure. However, there has been a lot of research in this area since Korean chatbots are in high demand and have great potential. Early chatbots used rules or simply retrieved stored responses. [5] However, these early systems could not handle the wide variety of conversations or understand the context well. To improve this, Vaswani *et al.* suggested neural-network model known as a Transformer. [4] Transformer uses a self-attention mechanism to understand relationships between word tokens, even when they are far apart in a sentence. This has led to promising results in many NLP tasks.

Despite these advancements, the computational overhead of Transformer models still remains a critical issue. As model sizes and data volumes increase, training and inference times become bottlenecks, adversely affecting the real-time interaction which is crucial for chatbot applications. To enhance the response speed of Korean chatbot models, this study proposes replacing LayerNorm [1] within Transformer architecture with the element-wise DynamicTanh (DyT) activation [6]. By employing DyT, it is anticipated that computational efficiency can be improved while maintaining the functionality of Layer Normalization, thereby enhancing the response speed of chatbots.

In this study, we train DyT-based and LayerNorm-based Transformer models using a Korean question-answer dataset given in the following Github: https://github.com/songys/Chatbot_data/

*<https://inspirehep.net/authors/2719371>

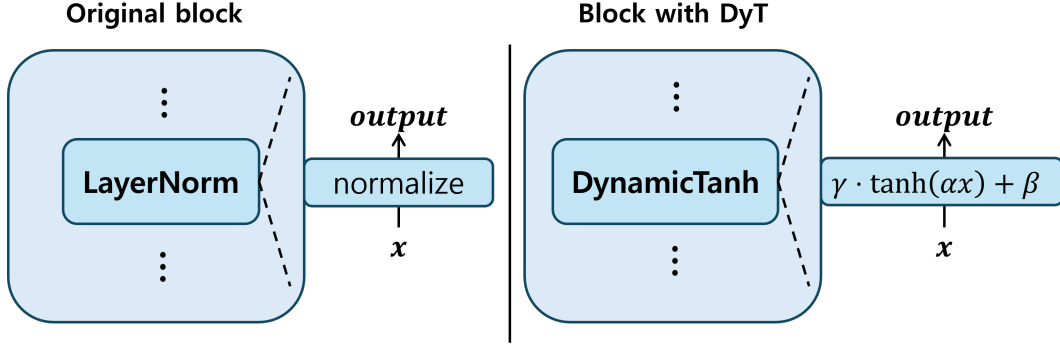


Figure 1: *Left*: original Transformer block with LayerNorm. *Right*: block with DyT layer. DyT can be substituted for the commonly used LayerNorm. Our experiments reveal that Transformer with DyT exceed the time efficiency of their LayerNorm counterpart.

40 ChatbotData.csv. Performances are evaluated using the BLEU score [2], with training/evaluation
 41 times measured to assess the efficiency of DyT. Experimental results demonstrate that the DyT
 42 model achieves comparable BLEU scores to the LayerNorm model while reducing training and
 43 evaluation times by approximately 15%. This finding suggests that DyT can be an effective method
 44 simultaneously improving both the performance and efficiency of Korean-language chatbots.

45 The remainder of this paper is organized as follows. Section 2 provides some preliminaries, intro-
 46 ducing Dynamic Tanh (DyT). Section 3 gives a detailed description of the proposed DyT-based
 47 Transformer model. Section 4 presents the experimental design and results. Section 5 discusses the
 48 analysis of the results, focusing on the time efficiency of DyT. Finally, Section 6 concludes the paper
 49 with a summary of our findings and potential directions for future research.

50 2 Preliminaries

51 This study proposes a method to improve the efficiency of the Transformer model by using DyT
 52 instead of LayerNorm. This chapter provides a short explanation of the Transformer model, Layer
 53 Normalization, and DynamicTanh, which are the foundations of the proposed method.

54 2.1 Transformers

55 The Transformer, proposed by Vaswani *et al.* in 2017 [4], has shown SOTA performance in sequence-
 56 to-sequence tasks, particularly in machine translation. The Transformer handles long-range dependen-
 57 cies within the input sequence using only a self-attention mechanism, without relying on recurrence
 58 neural networks (RNNs). This design facilitates parallel processing, significantly enhancing computa-
 59 tional efficiency compared to sequential models like seq2seq. [3]

60 2.1.1 Self-Attention

61 Self-attention is the core mechanism of the Transformer. Each token in the input sequence is trans-
 62 formed into Query, Key, and Value vectors. The attention weights between each pair of tokens are
 63 computed by taking the dot product of Query and Key, and then these weights are multiplied by the
 64 Value vectors to obtain a weighted sum, resulting in new representations for each token. In Multi-
 65 Head Attention, this process is performed in parallel across multiple "heads" to capture contextual
 66 information from various perspectives. The equation for scaled dot-product attention is as follows:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad (1)$$

67 where Q , K , and V are the Query, Key, and Value matrices, respectively, and d_k is the dimension of
 68 the Key vectors.

69 2.2 Layer Normalization (LayerNorm)

70 Layer Normalization (LayerNorm) [1] is a technique proposed to stabilize and accelerate the training
 71 of neural networks. LayerNorm normalizes the inputs within each training sample by computing the
 72 mean and variance across the neurons in a layer. The formulas are given as follows:

$$\begin{aligned} \text{Mean: } \mu &= \frac{1}{H} \sum_{i=1}^H x_i \\ \text{Variance: } \sigma^2 &= \frac{1}{H} \sum_{i=1}^H (x_i - \mu)^2 \\ \text{LayerNorm: } y_i &= \gamma \frac{(x_i - \mu)}{\sqrt{\sigma^2 + \epsilon}} + \beta \end{aligned}$$

73 where H is the number of neurons in the hidden layer, γ and β are learnable parameters, and ϵ is a
 74 small constant. LayerNorm does not depend on the batch dimension, making it effective for small
 75 batch sizes and even recurrent structures.

76 2.3 Dynamic Tanh (DyT)

77 DyT [6] is an operation designed as a replacement for LayerNorm layers in neural networks. The
 78 development of DyT is motivated by the observed similarity between the output patterns of LayerNorm
 79 layers and the shape of a scaled tanh function. Given an input tensor x , the DyT layer is defined by
 80 the following equation:

$$\text{DyT}(x) = \gamma \cdot \tanh(\alpha x) + \beta \quad (2)$$

81 In this formulation:

- 82 • α is a learnable scalar parameter. It plays a crucial role in enabling the input to be scaled
 83 differently depending on its range, thereby accommodating varying scales of x . This adaptive
 84 scaling is the reason behind the name "Dynamic" Tanh.
- 85 • γ and β are learnable, per-channel vector parameters, mirroring those used in standard
 86 normalization layers. These parameters serve the purpose of allowing the output to be
 87 rescaled to any desired scale.

88 .

89 3 Can DynamicTanh (DyT) replace LayerNorm?

90 This chapter presents the rationale for using Dynamic Tanh (DyT) as a replacement for LayerNorm
 91 in the Transformer model and briefly describes the model architectures used in the experiments. As
 92 proposed by [6], DyT exhibits a consistent similarity to the tanh shape, as observed in 2D image
 93 data-based self-supervised learning and diffusion models. (See figure 2 in [6].) Figure 3.1 shows the
 94 output distribution of LayerNorm after passing through several layers of the Transformer decoder
 95 for the Korean question-answering dataset. As can be seen in the figure, the output distribution of
 96 LayerNorm closely resembles the shape of a tanh function. Although most of the data distribution
 97 falls within the linear region, it can be observed that some data distribution exists in the non-linear
 98 regions of tanh. This observation suggests that the role of LayerNorm within the Transformer is to
 99 adjust the input values to a specific range and impart non-linearity. In particular, the similarity of the
 100 output distribution to the tanh function implies that an operation based on the tanh function could
 101 effectively replace the functionality of LayerNorm.

102 To verify the effectiveness of DyT, this study compares the performance of a baseline model using
 103 LayerNorm with a model using DyT under the same Transformer architecture and hyperparameter
 104 settings. The basic Transformer architecture of both models is the same as described in the previous

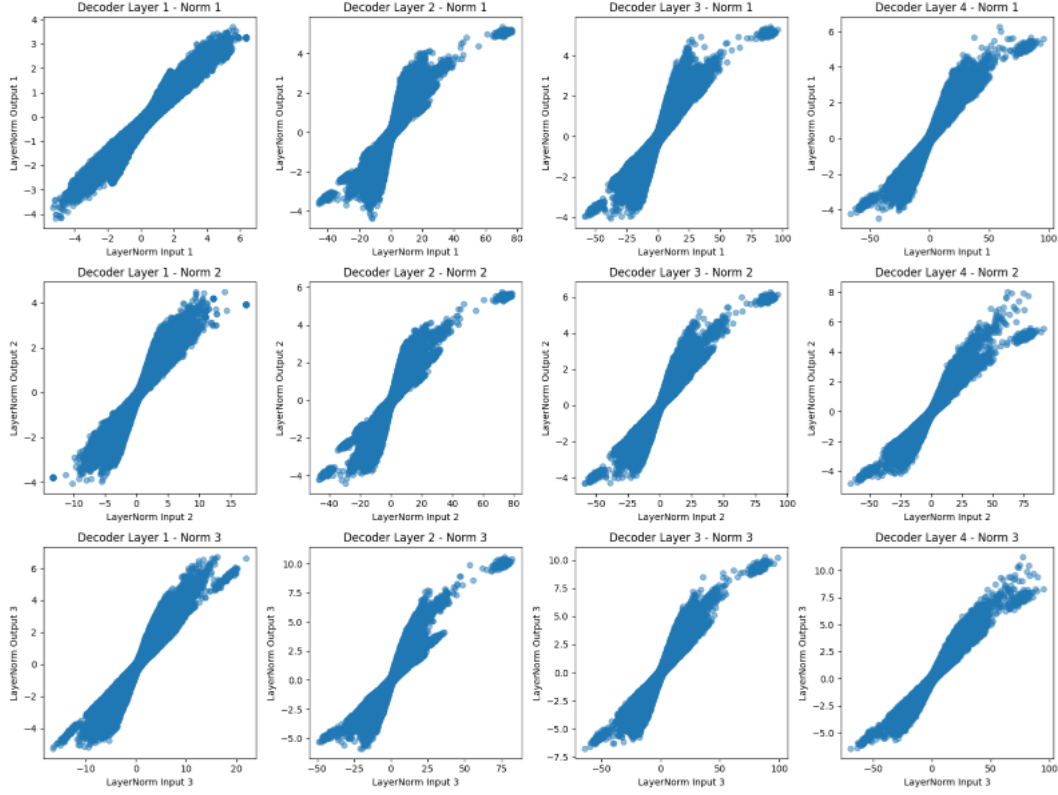


Figure 2: To illustrate the output characteristics of LayerNorm, we sampled a mini-batch and plotted the input against the output for four different LayerNorm layers. The resulting S-shaped curves strongly resemble the form of a tanh function (as shown in Figure 3). Notably, the more linear patterns observed in earlier layers can also be effectively represented by the central, near-linear portion of a tanh curve. This observation motivates our proposal of Dynamic Tanh (DyT) as a replacement in the Korean-text case.




Figure 3: The $\tanh(\alpha x)$ function with different α values.

chapter, with the only difference being the replacement of the internal LayerNorm layers with DyT layers. The next chapter will present the experimental results of training these models on the Korean question-answer dataset, followed by a comparative analysis of their BLEU scores and training/evaluation times.

4 Experiments

This section presents the experimental results of the LayerNorm-based Transformer model and the DyT-based Transformer model trained using a Korean question-answer dataset. Performance was

112 evaluated using the BLEU score, and model efficiency was compared by measuring training and
113 evaluation times.

114 4.1 Experimental Setup

115 4.1.1 Dataset

116 In this study, we used a publicly available Korean question-answering dataset (https://github.com/songys/Chatbot_data/ChatbotData.csv) for training.

118 4.1.2 Data Augmentation

119 To improve the generalization performance of the model, we applied data augmentation techniques.
120 Specifically, we used the following methods:

- 121 • **Lexical Substitution:** We replaced words based on their semantic similarity using a pre-
122 trained Korean text embedding model (FastText) provided by Facebook (<https://dl.fbaipublicfiles.com/fasttext/vectors-crawl/cc.ko.300.bin.gz>).
- 124 • **Random Insertion/Deletion:** We augmented the data by randomly inserting or deleting
125 words at arbitrary positions within the text. The words chosen for random insertion are those
126 with the highest cosine similarity.

127 4.1.3 Evaluation Metrics

128 The performance of the models was evaluated using the BLEU (Bilingual Evaluation Understudy)
129 score [2], a metric widely used in the field of NLP. The BLEU score measures the similarity between
130 the text generated by the model and the reference text written by humans.




Figure 4: Loss dynamics of LayerNorm-based and DynamicTanh-based models during training.

131

132 4.2 Results

133 We evaluated the performance and efficiency of LayerNorm and Dynamic Tanh (DyT) within our
134 Transformer model on a randomly selected subset (n=100) of the training data. The DyT model

Table 1: Comparison of average BLEU scores, execution times, and iterations per second for LayerNorm-based and DyT-based models. One can observe that the DyT model exceeds the time efficiency of the LayerNorm model.

Model	Average BLEU Score	Execution Time (seconds)	Iterations per second (it/s)
LayerNorm	64.70	88.88	6.27
DyT	64.64	74.47	8.03

achieved a comparable average BLEU score to the LayerNorm model, indicating successful replacement. However, DyT exhibited a significant advantage in terms of efficiency, demonstrating a 15% faster execution time and a greater number of minibatch iterations per second. The similar learning dynamics observed in Figure 4 further support the capability of DyT as a practical replacement. These findings suggest that DyT offers a computationally more efficient approach for Korean chatbot development using Transformers, without compromising performance.

5 Why is it More Efficient?

The experimental results demonstrate a remarkable observation: the DyT-based model exhibits an approximately 15% faster execution time compared to the LayerNorm-based model. This enhancement in time efficiency can be attributed to several key factors:

- LayerNorm involves statistical calculations for each input sample, including multiple binary operations. In contrast, DyT consists of a simple tanh operation. This difference in complexity may contribute to higher computational efficiency. The square root operation in LayerNorm, similar to division, is also computationally expensive. The element-wise independent operations of DyT are highly suitable for parallel processing, whereas statistical calculations can limit element-wise parallelism.
- The tanh function, a widely used activation function, has highly optimized implementations in deep learning frameworks. These optimizations can also contribute to faster operation for DyT.

6 Conclusion

This study explored replacing LayerNorm with DynamicTanh (DyT) in a Transformer-based Korean chatbot. Experiments on a Korean question-answer dataset showed DyT achieved comparable BLEU scores to LayerNorm, validating its performance-preserving capability. DyT reduced training and evaluation time by approximately 15%, indicating improved efficiency. These results suggest that DyT is a promising alternative to improve the efficiency of Transformer-based Korean chatbots without harming performance. Future research could explore the DyT performance across diverse Korean NLP tasks and larger datasets, and analyze its impact in various Transformer architectures. In conclusion, DyT shows significant potential as an efficient substitute for LayerNorm in Transformer-based Korean chatbot models.

References

- [1] Jimmy Lei Ba, Jamie Ryan Kiros, and Geoffrey E Hinton. Layer normalization. *arXiv preprint arXiv:1607.06450*, 2016.
- [2] Kishore Papineni, Salim Roukos, Todd Ward, and Wei-Jing Zhu. Bleu: a method for automatic evaluation of machine translation. In *Proceedings of the 40th annual meeting of the Association for Computational Linguistics*, pages 311–318, 2002.
- [3] Ilya Sutskever, Oriol Vinyals, and Quoc V Le. Sequence to sequence learning with neural networks. *Advances in neural information processing systems*, 27, 2014.
- [4] Ashish Vaswani, Noam Shazeer, Niki Parmar, Jakob Uszkoreit, Llion Jones, Aidan N Gomez, Łukasz Kaiser, and Illia Polosukhin. Attention is all you need. *Advances in neural information processing systems*, 30, 2017.

Table 2: Model Hyperparameters

Hyperparameter	Value
Total dataset size(augmented)	24770
Number of layers (n_layers)	4
Model dimension (d_model)	256
Number of attention heads (n_heads)	4
Feed-forward network dimension (d_ff)	1024
Dropout rate	0.5
Vocabulary size (vocab_size)	15000
Maximum positional length (pos_len)	50
Batch size	32
Number of epochs	20

- 175 [5] Joseph Weizenbaum. Eliza—a computer program for the study of natural language communica-
176 tion between man and machine. *Communications of the ACM*, 9(1):36–45, 1966.
- 177 [6] Jiachen Zhu, Xinlei Chen, Kaiming He, Yann LeCun, and Zhuang Liu. Transformers without
178 normalization. *arXiv preprint arXiv:2503.10622*, 2025.

179 Appendix

180 6.1 Hyperparameters

181 The hyperparameters used for the Transformer model in this study are detailed in Table 2.

183 6.2 Encoder LayerNorm Output




Figure 5: LayerNorm output of the encoder part.

184 Figure 5 illustrates the LayerNorm output of the encoder part. Most of the data points are located in
185 the linear region, which can be captured by the central part of a tanh curve.

186 6.3 Text Generation Examples

187 Table 3 provides examples of text generation results from the LayerNorm-based model and the
188 DyT-based model.

Table 3: Text Generation Examples

Example Sentence	LayerNorm Model Sequence	DyT Model Sequence
지루하다, 놀러가고 싶어.	재미 가 있을 거 예요	저도 놀고 싶어요
오늘 일찍 일어났더니 피곤하다.	시간이 좀 생길 거 예요	더 잠드는 건 어때요
간만에 여자친구랑 데이트 하기로 했어.	울고 있을 거 예요	잘 놀아요
집에 있다는 소리야.	집에 비밀은 시간이 예요	뭐라도 하세요
비가 와서 그런가 기분이 좀 별로야.	감성적이기 딱 좋죠	감성적이 예요